Transformer Architecture Overview

The Transformer is a deep learning architecture introduced in the paper Attention Is All You Need, published in 2

Self-Attention Mechanism

The self-attention mechanism allows the model to weigh the importance of different words in a sequence relative

Multi-Head Attention

Instead of performing a single attention function, the Transformer uses multi-head attention, which runs several a

Positional Encoding

Since the Transformer contains no recurrence or convolution, positional encodings are added to the input embed

Training and Applications

Transformers are trained using large-scale datasets with backpropagation and the Adam optimizer. They have b

The scalability of the Transformer architecture is one of its greatest strengths, as it parallelizes well across GPUs